

Diagnosing Conformal Prediction Failures Under Distribution Shift

A COVID-19 Case Study

Chorok Lee · KAIST
UAI 2026 — Amsterdam, Netherlands

KEY FINDING: SHAP Concentration predicts CP failure in GBMs ($\rho = 0.853$)

40% threshold achieves 87.5% LOO-CV accuracy | GBM models only

Problem & Motivation

Conformal Prediction (CP) provides distribution-free uncertainty, but **distribution shift breaks guarantees**.

Research Question:

Can we **predict** when CP fails **before** deployment?




Fig 1: COVID-19 distribution shifts (Feb-Jul 2020)

Key Insight

Models with **concentrated feature reliance** are vulnerable:

- Single feature drift $\Rightarrow$ large impact
- No redundant features for backup
- Calibration misalignment

High Concentration
 $\downarrow$
High Vulnerability

SALT Supply Chain Dataset

RelBench Supply Chain Benchmark

- 8 classification tasks
- COVID-19 temporal split
- Feb-Jul 2020 (onset $\rightarrow$ peak)

Task Type	Count
Sales-level	5
Item-level	3
Total SALT	8
+ External	+8

SHAP Concentration Method

Definition:

$$\text{SHAP Conc.} = \frac{\varphi_1}{\sum_{j=1}^p |\varphi_j|}$$

φ_1 = Top feature SHAP; $\sum |\varphi_j|$ = Total SHAP

Decision Threshold (40%):

<40% Robust
≥40% Vulnerable




Fig 2: Feature importance concentration

MAIN RESULTS

Strong Correlation: SHAP Concentration vs Coverage Drop

$$\rho = 0.853$$

Spearman Correlation

$p < 0.001$ | $n = 16$ tasks | 9 domains

SALT COVID (n=8): $\rho = 0.833$ ($p = 0.010$)




Fig 3: SHAP Concentration vs Coverage Drop (n=16)

Model Scope: GBM Only

Model	ρ	Works?
LightGBM	0.833	✓
CatBoost	0.667	✓
XGBoost	0.548	✓
Random Forest	0.30	✗
MLP	0.43	✗

GBM: concentrated-dependence
RF/MLP: global-sensitivity (different failure mode)

External Validation: WILDS + UCI

4 Benchmark Datasets

Dataset	Conc.	Cov.Δ	Pred.
CivilComments	26.2%	-0.7pp	Robust
Amazon	1.9%	+0.4pp	Robust
Covertypes	12.5%	+8.0pp	Robust
Adult	27.5%	+1.4pp	Robust


Conc.=SHAP Concentration, Cov.Δ=Coverage Change, Pred.=Prediction

91.7%

Combined Accuracy (11/12)

Practical Framework

Pre-Deployment Assessment



Conclusions & Future Work

Takeaways:

- SHAP Concentration predicts CP failures in GBMs: $\rho = 0.853$ ($p < 0.001$)
- 40% threshold: LOO-CV 87.5%; external validation 91.7%
- Scope: GBM classifiers; RF/MLP need different diagnostics
- Future Directions:
 - Mitigation: diversify features, add dropout/ensembles, trigger retraining
 - Prospective validation: recalibrate the 40% threshold on new domains
 - Streaming CP: monitor concentration and coverage drift online

Paper & Code:

github.com/ChorokLeeDev/
conformal-covid-uai2026

Contact:

chorok.lee@kaist.ac.kr



Scan for paper & code