
Diagnosing Conformal Prediction Failures Under Distribution Shift: A COVID-19 Case Study

Chorok Lee¹

¹Korea Advanced Institute of Science and Technology (KAIST), Daejeon, South Korea

¹choroklee@kaist.ac.kr

Abstract

We examine the reported association between feature-importance concentration and conformal coverage degradation in a COVID-era supply-chain case study and external datasets. Rounded historical summaries yield Spearman $\rho = 0.853$ for a selected 16-task subset and $\rho = 0.654$ for all 17 reported tasks. A statistical audit identified test-aware preprocessing, reuse of calibration labels for early stopping, different prediction-set conventions across cohorts, and missing numerical artifacts. These findings limit the original pre-deployment and replication claims. A recovered, separate follow-up panel of all eight SALT tasks (three fixed seeds per task) gives $\rho = -0.238$ under inclusive score inversion and -0.262 under crossing sets. Its frozen predictions and SHAP concentrations were independently regenerated from raw inputs and stored models, and all 24 follow-up fits were rerun. This panel does not reproduce the historical positive association; the original 50-seed models remain unavailable. We give a conditional score bound for an idealized probability mixture, whose weight is distinct from measured SHAP concentration, and withdraw unsupported class-count and ICC interpretations. The diagnostic and 40% threshold are unvalidated; no deployment rule is supported by these results.

1 INTRODUCTION

Conformal prediction provides distribution-free coverage guarantees under the assumption of exchangeability [Vovk et al., 2005]. However, real-world deployments face distribution shifts that violate this assumption. While prior work characterizes *how* conformal prediction degrades under shift [Tibshirani et al., 2019, Barber et al., 2023], a criti-

cal gap remains: **Can we identify which deployed models will fail before observing test data?**

We audit the historical proposal that **SHAP concentration**—the fraction of total feature importance concentrated in the top feature—could diagnose conformal prediction vulnerability before deployment. The original observational case study used a supply chain benchmark with temporal splits aligned to COVID-19 (February–July 2020), where 8 classification tasks experience common temporal periods yet exhibit coverage drops ranging from negligible ($<1\%$) to catastrophic ($>70\%$).

Our contributions:

1. **Exploratory diagnostic:** The reported SHAP concentration summaries correlate with coverage-drop summaries: $\rho = 0.853$ across the selected 16 tasks and $\rho = 0.654$ across all 17 tasks including Stack Overflow. These correlations can be recomputed from rounded published points, but the underlying experiment ledgers are missing and historical SALT/external set conventions differ. The recovered eight-task follow-up panel does not reproduce the positive association (Section 8). Prospective diagnostic value remains unestablished.
2. **Conditional theory:** Theorem 1 bounds an idealized probability mixture with explicit assumptions and a weight distinct from measured SHAP concentration. It does not verify those assumptions empirically or establish a guarantee for the archived experiments. Historical shift-detector comparisons are descriptive and await reproducible outputs.
3. **Exploratory framework:** A 40% concentration rule and protective-factor check summarize observed cases. They are retrospective hypotheses, with one external false negative and a protective rule derived from one false positive. Recovered single-seed retraining logs mix validation and test months; neither they nor the non-replicating follow-up panel validate a deployment recommendation.

2 RELATED WORK

Conformal Prediction. Vovk et al. [2005] introduced conformal prediction with exchangeability guarantees; Shafer and Vovk [2008] provide an accessible tutorial. Recent work extends to classification [Romano et al., 2020, Cauchois et al., 2021] and regression [Romano et al., 2019, Lei et al., 2018]. Angelopoulos and Bates [2023] provide a comprehensive survey. Ding et al. [2023] show that class-conditional coverage guarantees become harder to achieve as class count grows and propose clustered conformal prediction as a remedy; Gibbs et al. [2025] provide finite-sample conditional guarantees via a covariate-shift reweighting framework interpolating between marginal and conditional validity. Conditional and marginal guarantees must be distinguished: class count alone does not explain failure of marginal coverage. Historical partial-correlation summaries in Section 5.4 do not resolve the pipeline confounds.

Conformal Prediction Under Shift. Tibshirani et al. [2019] study covariate shift with known propensity scores. Podkopaev and Ramdas [2021] develop distribution-free uncertainty quantification for classification under label shift. Barber et al. [2023] provide a comprehensive treatment beyond exchangeability, bounding coverage loss by the total variation distance between test and calibration score distributions. Kasa and Taylor [2023] empirically characterize how CP degrades across many vision architectures and shift types, showing substantial degradation is common. We examine the historical hypothesis that source feature-importance concentration predicts tabular coverage degradation. The original test-aware fitting pipeline and the non-replicating recovered panel prevent claiming prospective predictive validity.

Adaptive Methods. Gibbs and Candès [2021, 2024] develop Adaptive Conformal Inference (ACI) for non-stationary settings, with extensions by Zaffran et al. [2022] for time series, Feldman et al. [2023] for online risk control, and Bhatnagar et al. [2023] for strongly adaptive online learning. Angelopoulos et al. [2024] generalize conformal prediction to broader risk measures. Kasa et al. [2025] propose entropy-based methods (ECP/EACP) that adapt prediction sets using unlabeled test data at deployment time—a related but distinct goal from the proposed source-only concentration summary. Computing that summary from validation covariates does not make the historically test-aware fitting pipeline a pre-deployment procedure. Historical clipped-ACI summaries report larger sets and improved coverage, without establishing transfer of the cited guarantee.

Shift Detection. WILDS [Koh et al., 2021] and Shifts [Malinin et al., 2021] provide benchmarks, while Gulrajani and Lopez-Paz [2021] show domain generalization methods often fail to improve over ERM under realistic distribution shifts. Gardner et al. [2023] benchmark tabular distribution shift across 15 tasks—the closest tabular-specific bench-

mark; our SALT dataset differs in providing a naturalistic COVID-19 temporal split with documented business-disruption context rather than curated benchmark partitions. Garg et al. [2022] use unlabeled test data to predict accuracy degradation—a related goal but requiring test-time observations. Kernel-based tests (MMD [Gretton et al., 2012]) and classifier two-sample tests (C2ST [Lopez-Paz and Oquab, 2017]) test distributional differences. The historical SALT summaries report weak severity correlations for these diagnostics, subject to the implementation and provenance defects in Section 6.3. They do not establish a general inability to predict failure severity or superiority of concentration. Miller et al. [2021] show that in-distribution (ID) validation accuracy predicts out-of-distribution (OOD) accuracy for image classifiers—a correlation that does not hold for conformal coverage in our tabular setting, where ID validation coverages cluster near the 90% target, spanning 83.6%–99.9% (Table 2), yet test-time coverage spans 12–100%, without establishing that concentration provides a validated improvement over accuracy-based assessments.

Interpretability for Reliability. We examine the proposed connection between SHAP [Lundberg and Lee, 2017, Lundberg et al., 2020] and reliability. Computing a concentration summary from source validation covariates is possible, but prospective failure prediction is an unvalidated hypothesis, not an established result of this study.

3 METHODOLOGY

3.1 DATASET AND TEMPORAL SPLITS

We use the SALT (Supply chain ALlocATion) dataset from RelBench [Robinson et al., 2024], a supply chain benchmark with temporal splits. All 8 classification tasks from the rel-salt benchmark are included; no tasks were excluded. External analyses report 9 additional datasets. We retain the historical $K \geq 4$ subset (8 external datasets plus 8 SALT tasks, $n = 16$) as an exploratory scope choice. APS mechanics do not rule out catastrophic coverage failure at $K = 2$ or 3, and no dated prospective exclusion plan is available. Including Stack Overflow ($K = 3$) yields the all-17 sensitivity result $\rho = 0.654$, $p = 0.004$; this sensitivity limits conclusions from the selected subset.

- **Train:** Before February 2020 (pre-COVID)
- **Validation:** February–July 2020 (COVID onset)
- **Test:** After July 2020 (COVID peak)

Audit of the historical data protocol. The supplied SALT code fits categorical encoders and the target vocabulary using train, validation, and test data. It also uses all validation labels for early stopping before dividing that sample into calibration and evaluation; the external script likewise reuses validation labels. Thus, although the reported SHAP

summaries use validation covariates, the full pipeline is not strictly pre-deployment and does not satisfy standard held-out split-conformal assumptions. The numerical tables below are historical reported results, not a validated rerun. Corrected evaluation requires train-only preprocessing, an independently defined class vocabulary, and separate tuning/calibration/evaluation samples. Threshold evaluation uses test-drop labels and is retrospective.

3.2 CONFORMAL PREDICTION SETUP

The archived experiments use two deterministic APS-related conventions. SALT set construction includes the class that crosses the calibrated cumulative-probability threshold. The external script instead computes coverage by the event $s(x, y) \leq \hat{q}$, while its size calculation includes the crossing class. These define different sets; the historical pooled endpoint and external coverage/size pairs therefore require reevaluation under one explicit convention. A new reference implementation uses the exact finite-sample order statistic and inverts a fixed-tie-order inclusive APS score. It has not generated the historical results reported here.

1. Train LightGBM classifier with fixed hyperparameters
2. Calibrate conformal predictor on 50% of validation set
3. Evaluate coverage on held-out validation and test sets

3.3 ENSEMBLE PROTOCOL

The historical design reports 50 seed runs (42–91) on the same data, with a fixed chronological validation split, t intervals across seeds and paired Wilcoxon comparisons. These summarize algorithmic variation conditional on one dataset and do not create 50 independent realizations of a distribution shift. Original seedwise results are unavailable for recomputing the intervals/tests.

3.4 FEATURE OVERLAP METRIC

Feature temporal stability uses Jaccard similarity:

$$J(x_j) = \frac{|A_{\text{train}} \cap A_{\text{test}}|}{|A_{\text{train}} \cup A_{\text{test}}|} \quad (1)$$

where A_{train} and A_{test} are the sets of unique values for feature x_j in train and test data. In the pre-deployment framework (Section 7), we substitute validation data for test data as a proxy.

3.5 SHAP CONCENTRATION DIAGNOSTIC

We define the SHAP concentration metric as:

$$C = \frac{\phi_1}{\sum_{j=1}^p \phi_j} \quad (2)$$

where ϕ_j is the mean absolute SHAP value of the j -th feature computed on validation data using TreeExplainer [Lundberg et al., 2020], and ϕ_1 is the importance of the top-ranked feature. For a fixed model, C can be computed from source validation covariates; the historical model-fitting pipeline nevertheless used test information.

3.6 WHY SHAP CONCENTRATION?

A natural question is whether simpler diagnostics suffice.

Feature overlap (Jaccard) is insufficient. Tasks relying on transaction IDs have Jaccard ≈ 0 while tasks with entity-based features have Jaccard > 0.5 (Table 3), yet Jaccard alone cannot explain the full range of coverage drops.

Information requirements. Target entropy changes require target covariates; target ECE additionally requires target labels. Concentration can use source covariates for a fixed model. The strongest association among the historical reported summaries does not establish superior prospective performance, particularly after metric selection and the failed follow-up replication (Section 5.4).

Why not raw feature importance? Raw importance values are not comparable across tasks (different scales, feature counts). Concentration normalizes to a $[0, 1]$ scale, enabling cross-task comparison.

Why top-1? Top-1 has the strongest observed association among the reported concentration summaries. The table reports eight candidate metrics (top-1/2/3/5/10, HHI, Gini, entropy). With the displayed SALT points, its unrounded two-sided Spearman $p = 0.0101755$ gives Holm-adjusted $p = 0.0814$ for this eight-test family (0.0509 even for the previously reported five-test family). Neither clears 0.05. The pooled $n = 16$ result reuses these eight SALT tasks and is not independent replication of metric selection. On the eight external rounded summaries alone, $\rho = 0.833$, $p = 0.0102$; this is descriptive arithmetic, pending provenance and pipeline reconciliation.

Model scope. The reported model-sensitivity correlations differ (LightGBM 0.833, CatBoost 0.667, XGBoost 0.548, RF 0.30, MLP 0.43). These observational comparisons do not identify architectural causes or establish that boosting necessarily reinforces single-feature dependence, bagging necessarily dilutes it, or SHAP approximation error explains the MLP result. TreeSHAP is exact for the specified tree-explanation model, but a causal mechanism would require controlled interventions and attribution-stability checks. Generalization beyond the evaluated models remains unverified.

Table 1: Per-Task Feature Importance Concentration. Top- k shows cumulative importance in the top k features (%). Drop = coverage degradation (pp). High Top-1 with catastrophic drop (e.g., s-payterms: 54.2% $\rightarrow$ 77.1 pp) versus high Top-2/3 with near-zero drop (s-office: 77.0%/84.8% $\rightarrow$ 0.1 pp) is an observed contrast, not evidence of a causal mechanism.

Task	Top-1	Top-2	Top-3	Drop
s-payterms	54.2	78.1	89.6	77.1
s-shipcond	50.7	77.5	92.5	71.6
i-shippoint	48.8	72.7	92.4	18.5
s-group	47.3	65.9	79.6	71.2
s-office	42.6	77.0	84.8	0.1
i-incoterms	28.9	42.3	55.1	11.3
i-plant	23.9	46.9	65.0	10.6
s-incoterms	23.7	47.2	69.1	8.5

Tasks sorted by Top-1 concentration. Top- k = $\sum_{j=1}^k \phi_j / \sum_{j=1}^p \phi_j$ (%). Drop = mean coverage drop (50 seeds, pp). Spearman ρ with drop: Top-1 = 0.833 ($p = 0.010$ before multiplicity correction); Top-2 = 0.524 ($p = 0.183$); Top-3 = 0.524 ($p = 0.183$); Top-5 = 0.690 ($p = 0.058$); Top-10 = 0.399 ($p = 0.328$). Alternative concentration metrics: HHI = 0.619 ($p = 0.102$); Gini = 0.500 ($p = 0.207$); Entropy = 0.571 ($p = 0.139$)—none significant. The contrast for sales-office motivates a mechanism hypothesis; the stability of individual features is not established by this correlation comparison.

4 THEORY: WHY CONCENTRATION RELATES TO FAILURE

We give a conditional score bound for an idealized probability mixture. Its weight λ is distinct from measured SHAP concentration C . The calculation is not a guarantee for the archived experiments and does not establish their mechanism.

Score and set convention. Fix a deterministic class-index tie order. Let $s(x, y^*)$ be the sum of probabilities through the true label in decreasing-probability order, and let $\mathcal{B}(x, y^*)$ contain the labels ranked after it. Then

$$s(x, y^*) = 1 - \sum_{y \in \mathcal{B}(x, y^*)} \hat{p}(y | x). \quad (3)$$

Each omitted probability is at most $\hat{p}(y^* | x)$, including ties. The score-inversion set is $\mathcal{C}_q(x) = \{y : s(x, y) \leq q\}$. This convention differs from the historical SALT crossing-label set; that distinction must be resolved before using this proposition to interpret its coverage.

Theorem 1 (Conditional bound for an idealized probability mixture). *Let $K \geq 2$, $\lambda \in [0, 1]$, and a fixed deterministic threshold $q \in [0, 1]$. Assume*

- (A1) $\hat{p}(y | x) = \lambda g(y | x_1) + (1 - \lambda)h(y | x_{\setminus 1})$, where g, h are probability distributions.
- (A2) At test time $g(y^* | x_1) \leq \varepsilon$ almost surely, for fixed $\varepsilon \in [0, 1]$.

- (A3) The joint distribution of $(h(\cdot | x_{\setminus 1}), y^*)$ is the same in calibration and test populations.

Write $\bar{h} = \mathbb{E}[h(y_{\text{cal}}^* | x_{\setminus 1}^{\text{cal}})]$. Then

$$s(x^{\text{test}}, y^*) \geq 1 - (K - 1)\{\lambda\varepsilon + (1 - \lambda)h(y^* | x_{\setminus 1}^{\text{test}})\}, \quad (4)$$

and consequently

$$\mathbb{E}[s_{\text{test}}] \geq 1 - (K - 1)\{\lambda\varepsilon + (1 - \lambda)\bar{h}\}. \quad (5)$$

Holding ε and $\bar{h}$ fixed, this lower bound is strictly increasing in λ if $\varepsilon < \bar{h}$. This concerns the bound, not monotonicity of the actual score expectation.

For $\lambda < 1$, score-inversion coverage obeys

$$\mathbb{P}(y^* \in \mathcal{C}_q(x^{\text{test}})) \leq \mathbb{P}(h(y_{\text{cal}}^*) \geq T(\lambda)), \quad (6)$$

where

$$T(\lambda) = \frac{(1 - q)/(K - 1) - \lambda\varepsilon}{1 - \lambda}.$$

Holding q, ε and the distribution of $h(y_{\text{cal}}^*)$ fixed, this upper bound is non-increasing when $\varepsilon \leq (1 - q)/(K - 1)$; the threshold is strictly increasing under strict inequality. At $q = 1, \varepsilon = 0$ it is constant. At $\lambda = 1$, if $q < 1 - (K - 1)\varepsilon$ coverage is zero; otherwise this argument yields only the trivial upper bound one.

Proof. The omitted sum in (3) contains at most $K - 1$ terms, each at most $\hat{p}(y^*)$. Substituting (A1)–(A2) gives (4); expectation and (A3) give (5). For the score-inversion event, $s \leq q$ implies the lower bound is at most q , yielding (6) when $\lambda < 1$. Its threshold derivative is $[(1 - q)/(K - 1) - \varepsilon]/(1 - \lambda)^2$. The $\lambda = 1$ case follows directly from the pointwise bound. $\square$

The proposition does not identify λ with the mean-absolute TreeSHAP ratio, whose usual LightGBM decomposition is in raw-margin space. Recalibrating a changed classifier also changes q and may change h , so the fixed-quantity monotonicity does not establish monotone empirical coverage across classifiers. The historical score comparisons in Appendix H are illustrative substitutions, not verification of (A1)–(A3).

5 RESULTS

5.1 COVERAGE DEGRADATION ACROSS TASKS

Despite experiencing common COVID-era temporal periods, the 8 supply chain tasks exhibit dramatically different coverage degradation (Table 2). The historical code uses a one-sided paired Wilcoxon alternative for positive drops and reports $p \leq 0.005$, with t intervals over seeds. Without the seed ledger these tests cannot be verified. Reported mean drops range from 0.1% to 77.1%.

Table 2: Historical Reported Coverage Degradation (50 Seeds; Original Runs Unverified). Reported 95% seed-level t intervals and paired Wilcoxon p -values are retained for the record; the missing seed ledger prevents verification.

Task	Cl	Val Coverage [95% CI]	Test Coverage [95% CI]	Drop [95% CI]	p	Cat
s-shipcond	45	93.5 [93.2, 93.7]	21.8 [14.1, 29.6]	71.6 [63.9, 79.4]	<0.001	SEV
s-group	459	83.6 [81.7, 85.5]	12.4 [3.1, 21.7]	71.2 [61.3, 81.0]	<0.001	SEV*
s-payterms	137	90.8 [90.7, 91.0]	13.7 [6.0, 21.5]	77.1 [69.4, 84.9]	<0.001	SEV
i-plant	35	92.0 [91.7, 92.3]	81.4 [79.0, 83.8]	10.6 [8.2, 13.0]	<0.001	ROB
i-shippoint	69	91.2 [90.9, 91.4]	72.7 [62.3, 83.0]	18.5 [8.2, 28.8]	0.005	ROB* [†]
s-incoterms	13	95.5 [95.3, 95.7]	87.0 [84.7, 89.3]	8.5 [6.2, 10.8]	<0.001	ROB
i-incoterms	13	95.0 [94.9, 95.2]	83.7 [80.9, 86.6]	11.3 [8.5, 14.1]	<0.001	ROB
s-office	25	99.9 [99.9, 100.0]	99.9 [99.9, 99.9]	0.1 [0.0, 0.1]	<0.001	ROB

SEV = Severe (>50% drop), ROB = Robust (<20% drop); *high model variance (CV > 50%), classified by median; s-group’s reported 83.6% validation coverage is below target; high class count alone does not explain failure of a marginal conformal guarantee. Calibration reuse, chronological shift and implementation require investigation; [†] meets At-risk criterion (drop >15 pp) under the decision framework (Table 11).

Table 3: Feature Overlap for Primary Features.

Task	Primary Feature	Jaccard	Drop
s-shipcond	SALESDOC (ID)	0.02	71.6%
s-group	SALESDOC (ID)	0.02	71.2%
i-incoterms	PRODUCT, PARTY	0.58	11.3%
s-office	SALESORG	0.61	0.1%

Drop = mean val – mean test coverage (50 seeds).

5.2 DIAGNOSTIC ANALYSIS

The historical summaries motivate two descriptive hypotheses; neither identifies the cause of the reported coverage variation:

Feature temporal stability. Table 3 reports larger drops for two low-overlap transaction-ID tasks than for two higher-overlap entity tasks. The original feature ledger is unavailable, and the recovered input-join defect prevents attributing this contrast to feature stability alone.

Feature importance concentration. Concentration is associated with the historical reported drops (Section 5.3), but does not establish vulnerability or a mechanism. The complete recovered follow-up panel does not reproduce the positive association.

5.3 SHAP CONCENTRATION CORRELATES WITH FAILURE

Across all 8 SALT multiclass tasks, SHAP concentration correlates with coverage degradation: Spearman $\rho = 0.833$, $p = 0.010$ (Table 4, Figure 2). A newly rerun percentile task bootstrap of the rounded summaries gives [0.29, 1.00] (10,000 resamples, seed 42), versus the historical report [0.30, 1.00]. Treating shared-domain tasks as exchangeable resampling units is not justified by this calculation; it is a sensitivity summary, not an established population in-

terval. A BCa interval has not been computed. The correlation magnitude is stable under leave-one-out analysis ($\rho \in [0.75, 0.96]$), though statistical significance is not maintained for 2 of 8 jackknife samples ($p = 0.052$) in the smaller $n = 7$ subsets. The original manuscript reported within-model SHAP bootstrap CVs below 1% and 95% CIs within ± 1 percentage point (1,000 resamples of 10,000 validation rows), but the original attribution/bootstrap ledger is unavailable; this stability claim has not been verified.

Historical mechanism hypothesis. Figure 1 illustrates the reported distinction. Catastrophic tasks: the dominant feature experiences a concentrated importance explosion ($4.5\times$ increase). Robust tasks: importance redistributes across features with ranking reshuffle ($10\text{--}20\times$ increases distributed). These historical plotted summaries have no recovered original attribution ledger and do not verify the probability-mixture assumptions of Theorem 1.

False positives. Under the mean-drop >15 pp criterion, sales-office is the SALT false positive; item-shippoint is at risk by mean although most seeds have small drops. The proposed protective rule was derived from sales-office alone, and inconsistent descriptions of its feature identity and overlap split require the original feature ledger. No deployment guarantee follows from this retrospective rule.

Dependence and uncertainty. SALT tasks share entities, periods and one benchmark. Their statistical independence is not established by between-task ICC. Seed-level intervals describe algorithmic variability conditional on these data; they do not quantify uncertainty across new datasets or future shifts. Negative validation/test agreement ICC cannot establish negative correlation or conservative paired tests (Appendix G).

Inferential status. Both the within-SALT and pooled cross-domain analyses are exploratory. The pooled endpoint includes the SALT tasks used in diagnostic development. Observed-effect power calculations do not confer confir-

Table 4: Stratified Correlation Analysis.

Group	n	ρ	Kendall τ	$p(\rho)$	Boot. 95% CI
Multiclass (SALT)	8	0.833	0.714	0.010	[0.30, 1.00]
Multiclass (4 dom.)	11	0.909 [†]	0.782	<0.001	[0.61, 1.00]
Multiclass (8 dom.)	15	0.882 [‡]	0.714	<0.001	[0.60, 0.97]
Multiclass (9 dom.)	16	0.853	0.667	<0.001	[0.50, 0.96]
Combined (10 dom.) [§]	17	0.654	—	0.004	—

LOO stability (SALT): $\rho \in [0.75, 0.96]$; 6/8 jackknife significant.

Multiclass (9 domains) is the primary result, excluding Stack Overflow under the historical scope choice. Under the at-risk criterion (drop > 15 pp), threshold performance at 40% is Precision = 0.83, Recall = 0.83 (Appendix D). [†]Single-seed external values; multi-seed-consistent value at $n = 11$ is

$\rho = 0.818$, $p = 0.002$ (see Appendix G). [‡] $n = 15$ is an intermediate analysis (7 external multiclass domains); $n = 16$ adds KDDCup99 as the 8th external domain to complete the 9-domain primary endpoint. Stack Overflow (3 classes, excluded scope subset) is excluded from all multiclass endpoints.

[§]Combined row includes all 17 tasks (8 SALT + 9 external, 10 domains: SALT + 9 external), adding Stack Overflow back to the $n = 16$ primary endpoint. Spearman $\rho = 0.654$, $p = 0.004$ (from Appendix G); Kendall τ not reported for this auxiliary endpoint.

matory status, and no dated prespecified analysis plan or raw prediction ledger is available in the released package.

Cross-domain extension. The pooled historical summaries yield $\rho = 0.853$ and Kendall $\tau = 0.667$ for the selected 16 tasks. The supplied eight external rounded points alone yield $\rho = 0.833$, $p = 0.0102$, while all 17 tasks yield $\rho = 0.654$, $p = 0.004$. These arithmetic checks do not validate the underlying missing runs, independence of task selection, or comparability of historical APS conventions. Covertypes is the reported catastrophic external case and KDDCup99 a low-concentration false negative.

5.4 DIAGNOSTIC COMPARISON

Table 5 preserves two alternative source-data diagnostic summaries. Their definitions can use a fixed model and source data, but the historical fitting pipeline used test information; these comparisons do not establish prospective performance.

Native feature importance concentration: top-1 LightGBM gain divided by total gain (zero SHAP cost, single seed; seed=42). The reported association is $\rho = 0.667$ ($p = 0.071$). Sales-office has 66.1% native concentration (vs. 42.6% SHAP) despite near-zero reported coverage drop. These summaries do not identify the cause of the difference between attribution metrics or establish a difference in predictive performance.

Ensemble disagreement: standard deviation of validation coverage across 50 seeds. Its reported association is $\rho = 0.452$ ($p = 0.260$), describing seed variability on the shared data without identifying a feature-level failure mechanism.

Table 5: Historical Source-Data Diagnostic Comparison ($n = 8$; Original Runs Unverified). Spearman ρ with 50-seed mean coverage drop. ρ_{partial} : partial Spearman controlling for the other variable in the SHAP-class-count pair.

Diagnostic	ρ	p	ρ_{partial}	p_{partial}
SHAP concentration	0.833	0.010	0.629	0.131
log(num_classes)	0.743	0.035	0.334	0.464
Native FI concentration	0.667	0.071	—	—
Ensemble disagreement	0.452	0.260	—	—

Class cardinality as a possible confound. The historical within-SALT partial correlations are non-significant (0.629, $p = 0.131$ for concentration; 0.334, $p = 0.464$ for log K). The reported pooled partial correlation is 0.771 ($p = 0.0008$) for concentration and -0.010 ($p = 0.97$) for log K . These are conditional associations, not evidence that concentration causally drives severity independently of model, domain, shift magnitude, or pipeline differences. Their numerical provenance and uncertainty require the original task/seed ledger.

We additionally compare against two *post-hoc* baselines computed from 10-seed ACI experiments:

Prediction entropy change: $\Delta H = H_{\text{test}} - H_{\text{val}}$, where $H = -\sum_c p_c \log p_c$ averaged over instances.

Expected Calibration Error change: $\Delta \text{ECE} = \text{ECE}_{\text{test}} - \text{ECE}_{\text{val}}$.

These distribution-comparison metrics require target observations, which can arrive during deployment. Concentration can be computed from source covariates, but its prospective risk value is not established, and the historical fitting pipeline used test information. The seven-task historical comparison is in Appendix F.

6 EXTENDED EXPERIMENTS

6.1 ADAPTIVE CONFORMAL INFERENCE ACROSS ALL TASKS

The historical clipped ACI variant reports coverage improvements on severe tasks with larger sets. It clips α_t to $[0.001, 0.999]$ and quantile levels to $[0, 1]$, so the original arbitrary-shift ACI guarantee does not automatically transfer. It assumes immediate labels in the processed test order. Missing prediction ledgers prevent rechecking the reported tradeoff, and a deployability judgment requires task-specific utility and feedback timing.




Figure 1: **Historical Feature Importance Dynamics (Original Runs Unverified)**. (A) Catastrophic task: dominant feature explodes $4.5\times$ while maintaining top rank (concentrated dependence). (B) Robust task: importance redistributes across features with ranking reshuffle (diversified dependence). (C–D) Importance dynamics and ranking stability. Despite $\sim 0\%$ Jaccard similarity, coverage drops range from negligible to catastrophic.

6.2 ALTERNATIVE CONFORMAL SCORING

The historical Regularized Adaptive Prediction Sets (RAPS) [Angelopoulos et al., 2021] summaries report task-dependent differences: substantially reduces drops for high-cardinality tasks (s-group: $73.5\% \rightarrow 10.4\%$; s-payterms: $79.9\% \rightarrow 35.1\%$) but not s-shipcond ($60.4\% \rightarrow 67.8\%$). These differences do not establish the idealized mechanism in Section 4; its mixture weight is not measured SHAP concentration (Appendix J; RAPS used random 50/50 calibration splits vs. deterministic main-experiment splits—within-experiment comparisons are unaffected).

6.3 COMPARISON WITH STANDARD SHIFT DETECTION

The historical summary reports MMD $p < 0.002$, C2ST accuracy $> 99.9\%$, PSI > 0.002 , and weak severity correlations on eight SALT tasks. Taskwise summaries and a parseable upstream script were recovered; the original supplementary copy does not parse. Recomputing correlations from these summaries reproduces their reported arithmetic, but does not reproduce the original detector fitting. The recovered PSI code discards observations outside training-bin endpoints and returns zero for a fully displaced uniform distribution or a constant binary feature changing from zero to one. Its historical severity comparison is therefore not a valid test of PSI. C2ST also uses its evaluation fold for early stopping. A nonzero PSI alone is not a significance test. These observations, if reproduced, would describe this suite; they do not establish that generic shift detectors cannot




Figure 2: **Historical Summary Association (Original Runs Unverified)**. Spearman $\rho = 0.853$, $p < 0.001$ across all 16 multiclass tasks in 9 domains (dark circles: 8 SALT supply-chain tasks; orange triangles: 8 external domains, 10-seed means). The plotted pre-deployment label and 40% cutoff are historical claims, not validated prospective performance; the recovered full follow-up panel does not reproduce this positive association. Within SALT alone: $\rho = 0.833$, $p = 0.010$ ($n = 8$).

predict severity in other settings.

6.4 RETRAINING ANALYSIS

Recovered single-seed shipping-condition logs identify the reported +18.9 pp as an equal-weight mean over eleven months, February–December 2020, pooling validation and test periods. It is not a test-only, per-record effect. From the rounded monthly log entries, July–December improvement averages 36.08 pp, with most of that change in November–December. Serial dependence, this denominator change, and the default entity-join defect preclude a general quarterly-retraining recommendation. The original $p = 0.036$ and three-task Holm-adjusted $p = 0.11$ remain unverified historical inference; the other two severe-task input ledgers are still unavailable.

6.5 CROSS-DOMAIN AND EXTERNAL VALIDATION

Historical external reports cover nine datasets, with the selected eight multiclass datasets contributing to the pooled association. Covertypes and Gas Sensor have supplied geographic/temporal loader definitions. The supplied KDD loader uses a positional split of one file rather than the claimed official train/test separation; Additional external/WILDS code has since been recovered, but code and manuscript protocols disagree (including Adult age 40 versus 50, CivilComments 200 versus 1,000 text features, and positional versus wilderness-based Covertypes splits). Original per-record results remain missing. Whether those con-

trols experience no relevant shift cannot be established from split names alone. These gaps prevent interpreting the pooled result as a validated cross-domain replication.

6.6 PLACEBO TEST

The historical pre-COVID placebo reports lower degradation for seven tasks, with ratios about 6–143 where the rounded denominator is positive; sales-office has a reported ratio of one. These summaries and asymmetric training windows do not isolate a causal COVID effect, and the newly recovered generating code does not restore the missing original result ledger.

7 DECISION FRAMEWORK

The original 40% rule and protective-factor exception are retained only as historical hypotheses. The recovered complete SALT panel has no task with a mean drop above 15 pp, despite concentrations above 40% for three tasks. It cannot estimate sensitivity to severe failures, and it does not support interpreting high concentration as an operational vulnerability label. Monitoring, retraining and abstention policies require separately specified utilities, feedback delays and prospective validation. No retraining deferral or deployment decision should follow from the historical threshold alone.

8 RECOVERED EVIDENCE AND FAILED REPLICATION

The owner’s separate research archive supplied pinned SALT tables, fitted models and per-record probability arrays for a later, explicitly different experiment. The fixed panel contains all eight tasks and seeds 42–44: 10,000 training rows, 3,000 distinct rows each for tuning, calibration and source audit, and 6,000 target rows per fit. Preprocessing dictionaries use sampled training only; the label alphabet uses full historical training. Model tuning and calibration are separated chronologically. Other task targets are excluded, while uppercase entity keys are retained by the declared feature contract. Temporal separation does not establish exchangeability or universal availability of every covariate at decision time.

The October audit restored the checksum-verified archive and independently reconstructed model inputs from raw entity IDs. All 24 stored models reproduced their calibration, audit and target probabilities exactly; native raw-margin TreeSHAP concentrations on the specified first 128 calibration rows also matched exactly. The audit then recomputed exact calibration order statistics and both set rules using fixed stable ties. A fixed-protocol refit of all 24 models also reproduced all 528 saved arrays and scientific metadata exactly. Model serialization differed only by an empty un-

used GPU-list field; every tree and other parameter matched. This is a current-snapshot reproduction, not recovery of the original 50-seed fitted models or their SHAP protocol.

Table 6: Recovered follow-up panel, three-seed task means. Concentration is native raw-margin TreeSHAP; drop is source-audit minus target coverage in percentage points. Both set rules use the same exact threshold.

Task	C (%)	Incl. drop	Cross. drop
item-incoterms	26.06	7.52	3.78
item-plant	22.56	2.63	2.28
item-shippoint	28.80	1.36	1.57
sales-group	52.10	3.70	3.65
sales-incoterms	30.86	1.61	0.47
sales-office	49.05	0.16	0.02
sales-payterms	60.01	1.82	1.86
sales-shipcond	32.41	2.24	2.26

Labels outside the historical-training alphabet are always counted as uncovered, including at an infinite threshold. Reported drops compare source-audit and target coverage; a small drop does not imply that target coverage attains 90%.

Eight-task Spearman correlations are -0.2381 for inclusive inversion and -0.2619 for crossing sets (nominal asymptotic $p = 0.5702$ and 0.5309). The existing exact-versus-linear quantile factorial replay yields the same task correlations under either threshold rule. These shared-domain task means are not 24 independent replications; the non-significant correlations do not establish equivalence or a population null. They show that this complete follow-up panel fails to reproduce the historical positive association.

A raw join audit also found that truncating entity tables at 1 July 2020, as the historical pipeline’s default `get_db()` call does, leaves every target entity unmatched on the pinned snapshot: 88,942/88,942 rows per sales task and 402,855/402,855 per item task. Training and validation joins have no missing IDs, and joining the full entity tables matches every target ID. This is a demonstrated input-contract artifact on the recovered snapshot. Missing original models, software/snapshot identities and prediction ledgers prevent attributing the published collapse, or its entire old/new discrepancy, to this artifact alone. An as-of feature contract is required before treating the full-table join as a prospective deployment pipeline.

9 CONCLUSION

The archived results suggest an association between concentrated feature importance and coverage degradation in the evaluated settings. The selected 16-task association is $\rho = 0.853$, compared with 0.654 when the excluded three-class task is retained. These are reproducible summary arithmetic, not an independently reproduced common-protocol

experiment. Test-aware preprocessing, calibration reuse, differing prediction-set conventions and missing numerical artifacts remain material limitations. The idealized theorem concerns a separate mixture weight under explicit assumptions and does not establish the proposed empirical mechanism. Controlled interventions, a leakage-free common pipeline, and independent prospective datasets are needed to assess mechanism, model scope and operational thresholds. The recovered complete follow-up panel fails to reproduce the positive association. The current revision therefore withdraws the operational diagnostic and unsupported hyperparameter-robustness claims while preserving original submitted artifacts for the record.

Author Contributions

C.L. conceived the research question, designed and executed all experiments, performed the statistical analysis, and wrote the paper.

Acknowledgements

We thank the RelBench team for providing the SALT dataset and the anonymous reviewers for their constructive feedback. The released supplementary package is partial; missing artifacts and detected implementation issues are documented in Appendix A.

References

- Anastasios N Angelopoulos and Stephen Bates. A Gentle Introduction to Conformal Prediction and Distribution-Free Uncertainty Quantification. *Foundations and Trends in Machine Learning*, 16(4):494–591, 2023.
- Anastasios N Angelopoulos, Stephen Bates, Jitendra Malik, and Michael I Jordan. Uncertainty Sets for Image Classifiers using Conformal Prediction. In *International Conference on Learning Representations*, 2021.
- Anastasios N Angelopoulos, Stephen Bates, Adam Fisch, Lihua Lei, and Tal Schuster. Conformal Risk Control. In *International Conference on Learning Representations*, 2024.
- Rina Foygel Barber, Emmanuel J Candès, Aaditya Ramdas, and Ryan J Tibshirani. Conformal Prediction Beyond Exchangeability. *The Annals of Statistics*, 51(2):816–845, 2023.
- Aadyot Bhatnagar, Huan Wang, Caiming Xiong, and Yu Bai. Improved Online Conformal Prediction via Strongly Adaptive Online Learning. In *International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pages 2337–2363. PMLR, 2023.

- Maxime Cauchois, Suyash Gupta, and John C Duchi. Knowing What You Know: Valid and Validated Confidence Sets in Multiclass and Multilabel Prediction. *Journal of Machine Learning Research*, 22(81):1–42, 2021.
- Tiffany Ding, Anastasios N Angelopoulos, Stephen Bates, Michael I Jordan, and Ryan J Tibshirani. Class-Conditional Conformal Prediction with Many Classes. In *Advances in Neural Information Processing Systems*, volume 36, 2023.
- Dheeru Dua and Casey Graff. UCI machine learning repository. University of California, Irvine, School of Information and Computer Sciences. <https://archive.ics.uci.edu>, 2017.
- Shai Feldman, Liran Ringel, Stephen Bates, and Yaniv Romano. Achieving Risk Control in Online Learning Settings. *Transactions on Machine Learning Research*, 2023.
- Josh Gardner, Zoran Popovic, and Ludwig Schmidt. TableShift: Benchmarking the robustness of tabular learning methods. In *Advances in Neural Information Processing Systems Datasets and Benchmarks Track*, 2023.
- Saurabh Garg, Sivaraman Balakrishnan, Zachary C Lipton, Behnam Neyshabur, and Hanie Sedghi. Leveraging Unlabeled Data to Predict Out-of-Distribution Performance. In *International Conference on Learning Representations*, 2022.
- Isaac Gibbs and Emmanuel Candès. Adaptive Conformal Inference Under Distribution Shift. In *Advances in Neural Information Processing Systems*, volume 34, pages 1660–1672, 2021.
- Isaac Gibbs and Emmanuel Candès. Conformal Inference for Online Prediction with Arbitrary Distribution Shifts. *Journal of Machine Learning Research*, 25(162):1–36, 2024.
- Isaac Gibbs, John J Cherian, and Emmanuel J Candès. Conformal Prediction with Conditional Guarantees. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 87(4):1100–1126, 2025.
- Arthur Gretton, Karsten M Borgwardt, Malte J Rasch, Bernhard Schölkopf, and Alexander Smola. A Kernel Two-Sample Test. *Journal of Machine Learning Research*, 13: 723–773, 2012.
- Ishaan Gulrajani and David Lopez-Paz. In Search of Lost Domain Generalization. In *International Conference on Learning Representations*, 2021.
- Kevin Kasa and Graham W. Taylor. Empirically Validating Conformal Prediction on Modern Vision Architectures Under Distribution Shift and Long-tailed Data. arXiv:2307.01088. Presented at ICML 2023 Workshop on Structured Probabilistic Inference & Generative Modeling, 2023.
- Kevin Kasa, Zhiyu Zhang, Heng Yang, and Graham W. Taylor. Adapting Prediction Sets to Distribution Shifts Without Labels. In *Proceedings of the Forty-First Conference on Uncertainty in Artificial Intelligence*, volume 286 of *Proceedings of Machine Learning Research*, pages 1990–2010. PMLR, 2025.
- Guolin Ke, Qi Meng, Thomas Finley, Taifeng Wang, Wei Chen, Weidong Ma, Qiwei Ye, and Tie-Yan Liu. LightGBM: A highly efficient gradient boosting decision tree. In *Advances in Neural Information Processing Systems*, volume 30, 2017.
- Pang Wei Koh, Shiori Sagawa, Henrik Marklund, Sang Michael Xie, Marvin Zhang, Akshay Balsubramani, Weihua Hu, Michihiro Yasunaga, Richard Lanus Phillips, Irena Gao, et al. WILDS: A benchmark of in-the-wild distribution shifts. In *International Conference on Machine Learning*, volume 139 of *Proceedings of Machine Learning Research*, pages 5637–5664. PMLR, 2021.
- Jing Lei, Max G’Sell, Alessandro Rinaldo, Ryan J Tibshirani, and Larry Wasserman. Distribution-Free Predictive Inference for Regression. *Journal of the American Statistical Association*, 113(523):1094–1111, 2018.
- David Lopez-Paz and Maxime Oquab. Revisiting Classifier Two-Sample Tests. In *International Conference on Learning Representations*, 2017.
- Scott M Lundberg and Su-In Lee. A Unified Approach to Interpreting Model Predictions. In *Advances in Neural Information Processing Systems*, volume 30, 2017.
- Scott M Lundberg, Gabriel Erion, Hugh Chen, Alex De-Grave, Jordan M Prutkin, Bala Nair, Ronit Katz, Jonathan Himmelfarb, Nisha Bansal, and Su-In Lee. From local explanations to global understanding with explainable AI for trees. *Nature Machine Intelligence*, 2(1):56–67, 2020.
- Andrey Malinin, Neil Band, Yarin Gal, Mark J. F. Gales, Alexander Ganshin, German Chesnokov, Alexey Noskov, Andrey Ploskonosov, Liudmila Prokhorenkova, Ivan Provilkov, Vatsal Raina, Vyas Raina, et al. Shifts: A dataset of real distributional shift across multiple large-scale tasks. In *Advances in Neural Information Processing Systems Datasets and Benchmarks Track*, 2021.
- John P. Miller, Rohan Taori, Aditi Raghunathan, Shiori Sagawa, Pang Wei Koh, Vaishaal Shankar, Percy Liang, Yair Carmon, and Ludwig Schmidt. Accuracy on the Line: on the Strong Correlation Between Out-of-Distribution and In-Distribution Generalization. In *International Conference on Machine Learning*, volume 139 of *Proceedings of Machine Learning Research*, pages 7721–7735. PMLR, 2021.

- Aleksandr Podkopaev and Aaditya Ramdas. Distribution-Free Uncertainty Quantification for Classification under Label Shift. In Cassio de Campos and Marloes Maathuis, editors, *Proceedings of the 37th Conference on Uncertainty in Artificial Intelligence*, volume 161 of *Proceedings of Machine Learning Research*, pages 844–853. PMLR, 2021. URL <https://proceedings.mlr.press/v161/podkopaev21a.html>.
- Joshua Robinson, Rishabh Ranjan, Weihua Hu, Kexin Huang, Jiaqi Han, Alejandro Dobles, Matthias Fey, Jan Eric Lenssen, Yiwen Yuan, Zecheng Zhang, Xinwei He, and Jure Leskovec. RelBench: A Benchmark for Deep Learning on Relational Databases. In *Advances in Neural Information Processing Systems Datasets and Benchmarks Track*, 2024.
- Yaniv Romano, Evan Patterson, and Emmanuel Candès. Conformalized Quantile Regression. In *Advances in Neural Information Processing Systems*, volume 32, 2019.
- Yaniv Romano, Matteo Sesia, and Emmanuel Candès. Classification with Valid and Adaptive Coverage. In *Advances in Neural Information Processing Systems*, volume 33, pages 3581–3591, 2020.
- Glenn Shafer and Vladimir Vovk. A Tutorial on Conformal Prediction. *Journal of Machine Learning Research*, 9: 371–421, 2008.
- Ryan J Tibshirani, Rina Foygel Barber, Emmanuel Candès, and Aaditya Ramdas. Conformal Prediction Under Covariate Shift. In *Advances in Neural Information Processing Systems*, volume 32, 2019.
- Vladimir Vovk, Alex Gammerman, and Glenn Shafer. *Algorithmic Learning in a Random World*. Springer, 2005.
- Margaux Zaffran, Olivier Féron, Yannig Goude, Julie Josse, and Aymeric Dieuleveut. Adaptive Conformal Predictions for Time Series. In *International Conference on Machine Learning*, volume 162 of *Proceedings of Machine Learning Research*, pages 25834–25866. PMLR, 2022.

Diagnosing Conformal Prediction Failures Under Distribution Shift: A COVID-19 Case Study (Supplementary Material)

Chorok Lee¹

¹Korea Advanced Institute of Science and Technology (KAIST), Daejeon, South Korea

¹choroklee@kaist.ac.kr

A REPRODUCIBILITY DETAILS

The released supplement is incomplete. It contains six Python scripts but no original prediction ledger, checkpoints, split manifests or numerical result files. The packaged shift-detector script has a syntax error, the SHAP script depends on an absent generator, and only four of the nine external loaders are supplied. The model-sensitivity, WILDS, retraining, placebo and several ablation analyses cannot currently be reproduced from this package. An account-wide recovery subsequently found most missing source programs, two historical summary files, selected monthly retraining logs, and a separate private follow-up archive. The original 50-seed model/prediction/SHAP ledgers remain unrecovered. Section 8 reports a newly executed replay of the follow-up raw inputs, models and arrays, with an aggregate-only reproducibility record in `audit/2026-10-08/recovery`. Original submissions and the private raw payload are preserved.

A.1 LIGHTGBM HYPERPARAMETERS

The archived 50-seed LightGBM [Ke et al., 2017] script specifies `objective=multiclass`, `boosting=gbd`, `num_leaves=31`, `learning_rate=0.05`, `feature_fraction=0.8`, `bagging_fraction=0.8`, `bagging_freq=5`, `num_boost_round=500`, `early_stopping_rounds=50`, seeds 42–91 (50 seeds), with no task-specific hyperparameter search. These settings do not describe every later follow-up or alternative-model run; original fitted-model provenance remains unavailable.

A.2 CONFORMAL PREDICTION

Historical calibration uses the first half of validation records and evaluation the second half, after the same validation labels have already been used for early stopping. This chronological split need not be exchangeable even within validation. Historical code uses NumPy interpolated quantiles at $\min(\lceil (n+1)(1-\alpha) \rceil / n, 1)$. The corrected reference instead selects the exact $\lceil (n+1)(1-\alpha) \rceil$ th order statistic, using $+\infty$ when the required rank exceeds n ; it also applies one score-inversion convention to coverage and set size. No historical results were replaced by this reference.

Calibration and test set sizes. Sales-level tasks (s-shipcond, s-payterms, s-group, s-incoterms, s-office): $n_{\text{val}} = 71,474$, $n_{\text{cal}} = 35,737$, $n_{\text{test}} = 88,942$. Item-level tasks (i-plant, i-shippoint, i-incoterms): $n_{\text{val}} = 293,896$, $n_{\text{cal}} = 146,948$, $n_{\text{test}} = 402,855$. At $\alpha = 0.1$, the historical quantile level $u_n = \lceil (n+1)(1-\alpha) \rceil / n$ exceeds $1-\alpha$ by 4.757×10^{-5} for sales tasks and 1.225×10^{-5} for item tasks. These are level differences, not score-threshold or coverage bounds. Atoms or gaps in the score distribution can amplify a small level change, and interpolation can differ from exact order-statistic calibration. The missing original scores prevent quantifying either effect on historical target coverage.

A.3 SHAP COMPUTATION

Method: TreeExplainer [Lundberg et al., 2020]. Subsample: 10,000 per split. Aggregation: mean absolute SHAP across classes. Concentration metric (Eq. 2): computed on validation data only.

A.4 BASELINE DIAGNOSTICS

Native feature importance concentration uses a single model (seed=42) with LightGBM’s gain-based importance; concentration is top-1 gain divided by total gain. Ensemble disagreement is the standard deviation of validation coverage across 50 seeds. Both are computed without SHAP.

A.5 COMPUTATIONAL RESOURCES

Standard CPU (8 cores, 8GB RAM). Wall-clock: ~ 3 –4 hours for full 8-task, 50-seed suite. Software: Python 3.9, LightGBM 3.3, SHAP 0.41.

A.6 EXTERNAL DATASET PROTOCOLS

Table 7 retains the reported external dataset protocols. The supplied code does not establish a common pipeline: SALT and external coverage use different set conventions, validation labels are reused for tuning, and recovered external loaders conflict with several historical descriptions while original per-record results remain unavailable. In particular, the supplied KDD loader contradicts the previously stated official-test protocol, as noted below. Historical results need provenance reconciliation and a common-protocol rerun.

Table 7: External Dataset Protocols. Historical reported protocols; supplied external code halves validation for calibration/evaluation; n_{tr} , n_{te} = approximate train/test sizes used; Seeds = 10 for all external datasets. Type: **DS** = documented real-world shift; **NC** = intended null-shift comparison (random/pre-defined split; absence of relevant shift and low concentration are not established by the split label).

Dataset	Shift mechanism	Type	K	n_{tr}	n_{te}	Source
Coverttype	Geographic: areas 1–2 (train), area 3 (val/cal), area 4 (test)	DS	7	290,680	36,968	[Dua and Graff, 2017]
Gas Sensor	Temporal batch drift: batches 1–6 (train), 7–8 (cal), 9–10 (test)	DS	6	9,840	4,070	[Dua and Graff, 2017]
KDDCup99	Supplied code: positional 60/20/20 of one 10% file; shift unverified	—	11	395,073	98,769	[Dua and Graff, 2017]
PAMAP2	Population shift: subject-based split (unseen subjects at test)	DS	12	1,674,379	268,493	[Dua and Graff, 2017]
Shuttle	Null-shift control: UCI Statlog; 60/20/20 random split	NC	7	46,400	11,600	[Dua and Graff, 2017]
Avila Bible	Null-shift control: copyist identification; 50/50 random split	NC	12	10,430	10,437	[Dua and Graff, 2017]
Pendigits	Null-shift control: UCI pre-defined split (digit recognition)	NC	10	7,494	3,498	[Dua and Graff, 2017]
Satimage	Null-shift control: satellite image; 60/20/20 random split	NC	6	5,144	1,286	[Dua and Graff, 2017]
Stack Overflow	Category shift (3-class subset); excluded from $n = 16$ historical selected endpoint	—	3	65,105	127,370	[Dua and Graff, 2017]

For Coverttype, the supplied loader uses wilderness areas 1–2 for training, 3 for validation and 4 for testing. These geographic indicators are not the cover-type labels (e.g., Spruce-Fir or Lodgepole Pine). For KDDCup99, the supplied loader fetches one 10% dataset, filters classes with at least 100 observations over the whole file, encodes object columns before splitting, and uses positional 60/20/20 partitions; no separate official test file is fetched. Historical table sizes and results need reconciliation with run provenance. The remaining reported splits are retained as historical descriptions where loaders are absent; SHAP uses the validation pool, which the external code also uses for early stopping before calibration/evaluation.

Table 8: Placebo Test: Pre-COVID vs COVID Degradation (50-Seed Mean Drops).

Task	Placebo	COVID	Ratio
s-shipcond	0.5%	71.6%	143×
s-group	2.0%	71.2%	36×
s-payterms	0.0%	77.1%	>100×
i-plant	1.8%	10.6%	6×
i-shippoint	1.5%	18.5%	12×
s-incoterms	1.2%	8.5%	7×
i-incoterms	0.6%	11.3%	19×
s-office	0.1%	0.1%	1×

B PLACEBO TEST DETAILS

C HIGH MODEL VARIANCE ANALYSIS

Tasks marked with * in Table 2 exhibit extreme model variance ($CV > 50\%$):

- **s-group**: Mean 12.4%, median **0.5%**. Most models fail; outliers inflate mean.
- **i-shippoint**: Mean 72.7%, median **89.9%**. Most models succeed; failures deflate mean.

D THRESHOLD SENSITIVITY ANALYSIS

Table 9 shows threshold performance across $n = 16$ multiclass tasks. Note that the 40% threshold is derived from the SALT concentration gap (in-sample); the eight external datasets were not used to fit the displayed cutoff, but their post-outcome scope selection and test-aware pipelines prevent a clean prospective validation interpretation. On the external subset alone at 40%: Precision = 1.00, Recall = 0.50 (1 TP: Covertypes; 1 FN: KDDCup99; 6 TN). Low external recall reflects sparse catastrophic cases; the original primary claim rested on the $n = 16$ correlation ($\rho = 0.853$) rather than threshold classification. Ground-truth labels use the at-risk rule (mean coverage drop >15 pp, consistent with 50-seed means in Table 2). For high-variance tasks, the mean and median labels can differ: i-shippoint has mean drop 18.5 pp (At-risk) but median drop 1.2 pp (robust); Table 2 labels it ROB by median for descriptive purposes, while the framework uses the mean-based threshold for threshold evaluation.

Table 9: Threshold Sensitivity for SHAP Concentration Across 16 Multiclass Tasks (8 SALT + 8 external; Step 2 only, no protective-factor check). At-risk defined as coverage drop >15%.

Threshold	Prec	Rec	F1	TP	FP	FN
25%	0.63	0.83	0.71	5	3	1
30%	0.71	0.83	0.77	5	2	1
35%	0.83	0.83	0.83	5	1	1
40%	0.83	0.83	0.83	5	1	1
45%	1.00	0.83	0.91	5	0	1
50%	1.00	0.33	0.50	2	0	4

At 40%, the single FP is s-office (concentration 42.6%) for which the original paper proposed a protective-feature explanation (reported SALESORG Jaccard = 0.61, original feature ledger unavailable), and the single FN is KDDCup99 (intermediate regime; $C = 21.1\%$, drop = 15.9 pp). At 30%, the two FPs are s-office and Shuttle ($C = 30.7\%$); at 25%, i-incoterms ($C = 28.9\%$) is additionally flagged. Retrospectively exempting s-office raises the displayed precision to 1.00 while leaving recall 0.83; this exception was derived from the same case and is not independently validated.

D.1 LEAVE-ONE-OUT CROSS-VALIDATION

The original paper reported leave-one-task-out threshold calculations on SALT ($n = 8$). Reusing the task collection used to select the diagnostic does not provide an independent prospective test. For each fold, we exclude one task, compute the

optimal threshold from the remaining 7 tasks, and predict the held-out task’s vulnerability status.

Summary. LOO-CV achieves 87.5% accuracy (7/8 correct predictions). The threshold remains stable across folds: $43.1\% \pm 5.0\%$ (mean $\pm$ std). The effect size separating failed from robust tasks is Cohen’s $d = 3.08$ (large), with bootstrap 95% CI [0.31, 1.00] for the underlying correlation.

Table 10: LOO-CV Per-Fold Predictions (SALT, $n = 8$). Threshold = optimal threshold from 7 training tasks; Pred = predicted status at that threshold; Actual = observed outcome (Failed if drop > 15 pp).

Task	Conc. (%)	Drop (pp)	Thresh (%)	Pred	Actual	Correct
s-shipcond	50.7	71.6	45	Vuln	Failed	✓
s-group	47.3	71.2	45	Vuln	Failed	✓
s-payterms	54.2	77.1	45	Vuln	Failed	✓
i-plant	23.9	10.6	45	Robust	Robust	✓
i-shippoint	48.8	18.5	45	Vuln	Failed	✓
s-incoterms	23.7	8.5	45	Robust	Robust	✓
i-incoterms	28.9	11.3	45	Robust	Robust	✓
s-office	42.6	0.1	30	Vuln	Robust	×

Effect size analysis. Failed tasks ($n = 4$) have mean concentration $50.2\% \pm 2.6\%$, while robust tasks ($n = 4$) have mean concentration $29.8\% \pm 7.7\%$. The separation of 20.4 pp yields Cohen’s $d = 3.08$ (large effect), describing separation in these retrospectively labeled eight tasks; it does not establish prospective reliability.

Misclassification analysis. Sales-office is a false positive under the concentration-only rule. Its stable-feature explanation is a hypothesis derived from this single case. The historical text inconsistently describes SALESORGANIZATION as dominant or secondary and its Jaccard overlap as train-test or train-validation; the original feature ledger is required to resolve this. A deployable protective-factor rule must use information available before test outcomes. High concentration is not necessary for failure, as the reported low-concentration KDD false negative demonstrates.

E DECISION FRAMEWORK VALIDATION

Historical seed-stability proposal (withdrawn). The original paper recommended at least five seeds and a vulnerability override when a concentration confidence bound crossed 40%. These operational suggestions were not prospectively validated and are withdrawn. Seed variation estimates conditional fitting variability; it does not turn the concentration cutoff into a calibrated risk rule. The reported KDDCup99 concentration $21.1 \pm 7.5\%$ and SALT within-task CV claims require the missing original seed ledger for verification.

Threshold selection. The historical 40% threshold was proposed based on the empirical distribution of SHAP concentration values on validation data: a natural gap separates low-concentration tasks (24–29%) from high-concentration tasks (43–54%), visible in Table 11. We emphasize that this threshold is *exploratory*, derived from $n = 8$ multiclass tasks. The active revision withdraws its operational interpretation; independent prospective validation would be necessary for any renewed proposal.

Historical cross-domain classification summary. The 40% threshold applied without tuning to 9 external domains yields 7/9 deterministic or near-deterministic outcomes (6 at 10/10, Shuttle at 9/10), with Covertypes correctly flagged as catastrophic and Gas Sensor correctly unflagged as robust. KDDCup99 is the principal intermediate-regime false negative, while Stack Overflow (3 classes) is a false positive excluded from the selected endpoint. Across the $n = 16$ multiclass set: $\rho = 0.853$, $p < 0.001$ (Section 6.5). The sensitivity analysis in Appendix D reports threshold performance across 30–50%; sparse catastrophic external cases and the failed full-panel follow-up do not support operational use of this threshold.

F HISTORICAL BASELINE COMPARISONS AND INFORMATION REQUIREMENTS

Concentration can be computed on source covariates for a fixed model, whereas target entropy requires target covariates and target ECE requires target labels. The historical fitting pipeline nevertheless used test information. This section preserves reported entropy and ECE changes without establishing prospective diagnostic performance.

The historical summaries report decreasing entropy on the three tasks with coverage drops above 70% and increasing entropy on two tasks with smaller drops. These associations do not identify concentrated dependence as the cause, nor do they

Table 11: Historical Concentration Classification (Not a Deployment Rule) Applied to All Multiclass Tasks (Actual = At-risk if drop > 15 pp).

Task	Conc. (%)	Step 2	Actual (> 15 pp)	Note
<i>SALT (in-sample)</i>				
s-payterms	54.2	VULN	At-risk	—
s-shipcond	50.7	VULN	At-risk	—
i-shippoint	48.8	VULN	At-risk*	High variance [†]
s-group	47.3	VULN	At-risk*	—
s-office	42.6	VULN	ROB	Protective [‡]
i-incoterms	28.9	ROB	ROB	—
i-plant	23.9	ROB	ROB	—
s-incoterms	23.7	ROB	ROB	—
<i>External (held-out)</i>				
Covertypes	49.8	VULN	At-risk	10/10 seeds
Shuttle	30.7	ROB	ROB	9/10 seeds
Avila Bible	20.5	ROB	ROB	10/10 seeds
PAMAP2	19.2	ROB	ROB	10/10 seeds
KDDCup99	21.1	ROB	At-risk [§]	Seed-dependent (FN)
Pendigits	14.5	ROB	ROB	10/10 seeds
Satimage	9.0	ROB	ROB	10/10 seeds
Gas Sensor	7.3	ROB	ROB	10/10 seeds
Stack Overflow	48.9	VULN	ROB	Excluded-scope false positive [¶]

[†]Median drop 1.2%; most seeds maintain coverage. [‡]SALESORGANIZATION: train-validation Jaccard=0.61, importance>15%.

[§]Concentration below threshold ($C = 21.1 \pm 7.5\%$), but mean drop exceeds the at-risk cutoff (15.9 pp); outcome seed-dependent (range −0.8 to 73.5 pp). In-sample F1 = 0.89 (1 FP: s-office). External: Covertypes deterministic catastrophic (10/10); five low-concentration datasets deterministic robust (10/10); Shuttle near-deterministic (9/10); KDDCup99 is the intermediate-regime false negative. [¶]Stack Overflow ($K = 3$, $C = 48.9\%$) is a reported false positive with drop −7.0 pp. Its exclusion from the historical $n = 16$ subset is not justified by a general APS ceiling: binary and three-class predictors can also lose all coverage under shift. The all-17 association is $\rho = 0.654$, $p = 0.004$.

Table 12: Historical Source-Data (SHAP) and Target-Data (Entropy, ECE) Diagnostic Summaries (10-Seed ACI Experiments; Original Runs Unverified). SHAP concentration uses validation covariates for a historically test-aware fitted model. Δ Entropy and Δ ECE require test-time observations.

Task	Cat	SHAP Conc. [†]	Δ Entropy	Δ ECE
<i>Severe tasks</i>				
s-shipcond	SEV	50.7%	−0.746	+0.408
s-group	SEV	47.3%	−1.274	+0.176
s-payterms	SEV	54.2%	−1.979	+0.629
<i>Robust tasks</i>				
i-plant	ROB	23.9%	+0.181	+0.131
i-shippoint	ROB	48.8%	+0.469	+0.073
s-incoterms	ROB	23.7%	+0.364	+0.350
s-office	ROB	42.6%	+0.032	+0.104

[†]Computed on validation covariates for a historically test-aware fitted model; not a validated pre-deployment procedure. item-incoterms omitted (3-seed pilot only).

establish that an entropy-based detector would miss these cases under a specified alarm rule. ECE increases in the reported summaries likewise do not establish a general inability to distinguish severity. The original prediction ledger and a common correctly joined pipeline are needed to compare these diagnostics; no prospective superiority claim is made.

G TASK INDEPENDENCE ANALYSIS

We analyze pseudo-replication in the 50-seed design using intraclass correlation (ICC).

Between-task ICC. The archived summary reports $\text{ICC} = 0.675$ for coverage drops and 0.720 for test coverage. These describe a fitted variance decomposition across task/seed levels; they do not prove that tasks sharing one benchmark are independent. A design-effect conversion of 400 seed/task measurements cannot increase the number of independently observed tasks beyond the eight available tasks.

Per-task ICC(val,test). The previous analysis interpreted negative one-way agreement ICC as negative seedwise correlation and converted 50 pairs into effective sample sizes of 55–250. That inference is invalid. For example, when test coverage equals validation coverage minus a constant, Pearson correlation is one while the supplied agreement ICC can be strongly negative. We therefore withdraw the effective-sample-size table and the assertion that paired Wilcoxon tests must be conservative. Appropriate paired inference uses the distribution of seedwise differences and remains conditional on the common data.

Multiple comparisons. The reported concentration search includes eight metrics. Using the unrounded SALT p -value, the smallest Holm-adjusted value is $8 \times 0.0101755 = 0.0814$. The original five-metric family would give 0.0509 , not a value below 0.05 . Intermediate and pooled subset analyses are retained as exploratory; reusing discovery tasks and selecting endpoints cannot be repaired solely by correcting three subset p -values.

Exact small-sample sensitivity. The 0.0101755 value above is SciPy’s asymptotic Spearman approximation. Enumerating all $8! = 40,320$ no-tie rank permutations for the displayed SALT $\rho = 5/6$ gives 620 outcomes at least as extreme in absolute correlation: two-sided $p = 0.015377$. The eight-test Holm first-step sensitivity is $8p = 0.1230$; the full eight-metric exact analysis has not been rerun. This calculation assumes exchangeability of task pairs, which is not established for related SALT tasks. It checks small-sample calibration of the reported summary, not the original experiments or population generalization.

Multi-seed aggregation. The primary $n = 16$ correlation ($\rho = 0.853$) uses 10-seed means for all external datasets and 50-seed means for SALT tasks. The early $n = 11$ result ($\rho = 0.909$) used single-seed external values; switching to multi-seed means at $n = 11$ yields $\rho = 0.818$, $p = 0.002$, with the reduction driven by KDDCup99 whose multi-seed mean drop (15.9 pp) exceeds its single-seed value (-0.8 pp). This intermediate-concentration regime ($C = 21.1 \pm 7.5\%$) exhibits high outcome variance characteristic of the boundary between robust and vulnerable. Extending to $n = 16$ with consistent multi-seed means yields $\rho = 0.853$ ($p < 0.001$). Including Stack Overflow (3 classes, excluded scope subset) to make $n = 17$ weakens correlation to $\rho = 0.654$.

H THEORY DETAILS AND SCORE CDFS

Scope of the proof. The proof in Section 4 uses a fixed deterministic rank order and a probability-mixture weight λ . It does not assume that classes tied with the true label are strictly lower in probability, and it handles $\lambda = 1$ separately. Its coverage statement is for the explicit score-inversion set with fixed q , not a claim about recalibrated models or the archived crossing-label implementation.

Illustrative bound calculation, not empirical verification. The historical comparison substituted mixture weight $\lambda = C$, $\varepsilon = 0$ and $\bar{h} = 1/K$, obtaining numbers below five reported mean scores. None of these substitutions verifies the probability-mixture assumptions. In particular the lower bound decreases with $\bar{h}$, so $1/K$ is not conservative for every residual model whose true-class probability is at least uniform. These comparisons do not validate a theorem for measured TreeSHAP concentration; empirical verification requires an identified residual decomposition and a consistent score/set implementation.

I APS SCORE DISTRIBUTION COMPARISONS

Table 13 reports two-sample KS statistics comparing calibration and test APS scores (seed=42, non-randomized APS). An inclusive APS score is cumulative probability through the true label, not a monotone transform of the true-label probability.




Figure 3: **Conformity score CDFs (seed=42)**. The archived curves suggest rightward score shifts for some high-drop tasks and different patterns for lower-drop tasks. No stochastic-dominance test or original score arrays are available.

For example, if the true label stays top-ranked, increasing its probability from 0.6 to 0.9 also increases its APS score from 0.6 to 0.9. Score shifts alone therefore do not identify changes in true-class confidence.

Table 13: KS Tests: Calibration vs. Test Conformity Score Distributions (Seed=42).

Task	KS	p	Mean _{cal}	Mean _{test}
<i>Catastrophic tasks (test scores > calibration)</i>				
s-shipcond	0.956	$< 10^{-10}$	0.54	0.98
s-payterms	0.748	$< 10^{-10}$	0.65	0.92
s-group	0.676	$< 10^{-10}$	0.98	1.00
<i>Robust tasks</i>				
i-plant	0.741	$< 10^{-10}$	0.73	0.86
s-incoterms	0.633	$< 10^{-10}$	0.75	0.87
s-office	0.994	$< 10^{-10}$	1.00	1.00
i-incoterms	0.559	$< 10^{-10}$	0.78	0.66
i-shippoint	0.505	$< 10^{-10}$	0.71	0.58

All KS tests are two-sided. Historical deterministic APS scores; a conservative split-conformal guarantee requires the held-out/exchangeability assumptions absent from the supplied protocol. For catastrophic tasks, test scores reportedly shift rightward; two-sided KS and mean shifts alone do not establish stochastic dominance or verify Theorem 1. Two lower-drop tasks (i-incoterms, i-shippoint) have lower reported mean test scores. This does not establish greater true-class confidence or a product-diversity mechanism.

J RAPS COMPARISON

Table 14 compares APS and RAPS coverage drops across all 8 tasks with 10-seed means ($\lambda = 0.01$, $k_{\text{reg}} = 5$).

Table 14: APS vs. RAPS Coverage Drop (10-Seed Means $\pm$ Std). $\lambda = 0.01$, $k_{\text{reg}} = 5$.

Task	Cl	APS Drop	RAPS Drop	Δ	Note
<i>Severe tasks</i>					
s-shipcond	45	60.4 ± 31.8	67.8 ± 25.4	+7.4	No improvement
s-group	459	73.5 ± 33.9	10.4 ± 1.3	-63.1	RAPS protects
s-payterms	137	79.9 ± 29.4	35.1 ± 28.6	-44.8	Partial recovery
<i>Non-severe tasks</i>					
i-plant	35	11.8 ± 9.0	13.1 ± 5.8	+1.3	Near-identical
i-shippoint	69	9.3 ± 27.6	20.5 ± 21.6	+11.2	RAPS worse
i-incoterms	13	10.8 ± 9.1	10.8 ± 9.1	0.0	Identical
s-incoterms	13	6.7 ± 6.8	6.7 ± 6.8	0.0	Identical
s-office	25	0.04 ± 0.03	0.04 ± 0.03	0.0	Identical

All values are 10-seed means $\pm$ std. APS drops differ from 50-seed Table 2 due to smaller seed range (10 vs. 50) and high variance. The key comparison is APS vs. RAPS within the same seed set. Note: RAPS experiments used a random 50/50 calibration split rather than the deterministic first-half/second-half split used in the main APS experiments (Appendix A); this protocol difference does not affect the within-experiment APS vs. RAPS comparison but means RAPS coverage values are not directly comparable to Table 2. The historical task differences do not identify whether class accumulation, input joins, score conventions or other model effects cause the apparent RAPS benefit. The original per-record comparison is unavailable. For item-shippoint ($K = 69$), the historical summary reports a larger RAPS coverage drop by 11.2 pp (paired Wilcoxon $p = 0.064$, 10 seeds). Missing paired outputs prevent verifying this test or attributing the difference to a specific calibration mechanism.

K MODEL CLASS SENSITIVITY

Table 15 preserves the reported seed-42 analyses with `RandomForestClassifier` (500 trees), `XGBClassifier`, and `CatBoostClassifier` (ordered boosting). Its `LightGBM` coverage-drop column repeats the historical 50-seed means from Table 2; it is not a matched seed-42 comparison across all four models. Shared historical preprocessing defects and missing original fitted-model provenance further limit interpretation of model-family differences.

Table 15: Historical SHAP Concentration and Coverage Drop Across Model Classes. LightGBM drops are 50-seed means; alternative-model drops are reported seed-42 values. Original runs are unverified.

Task	LGB	RF	XGB	CB	LGB	RF	XGB	CB
	Concentration (%)				Coverage Drop (%)			
s-shipcond	50.7	57.4	44.6	53.7	71.6	15.8	47.2	1.1
s-group	47.3	22.5	33.3	31.8	71.2	50.0	56.8	14.5
s-payterms	54.2	27.9	30.6	25.7	77.1	13.8	20.4	-5.3
i-shippoint	48.8	52.1	41.0	43.1	18.5	22.7	80.4	61.8
i-plant	23.9	26.8	26.0	29.0	10.6	15.3	9.7	8.5
s-incoterms	23.7	26.8	21.6	28.2	8.5	2.5	3.7	-1.5
i-incoterms	28.9	17.5	18.4	21.7	11.3	2.9	1.1	-0.4
s-office	42.6	24.5	36.8	33.5	0.1	0.0	0.0	-0.0

Historical reported results. The concentration-drop association varies by model class (Table 16). All three gradient-boosted implementations show positive correlation: LightGBM ($\rho = 0.833$, $p = 0.010$), CatBoost ($\rho = 0.667$, $p = 0.071$), and XGBoost ($\rho = 0.548$, $p = 0.16$). RandomForest ($\rho = 0.30$, $p = 0.47$) does not replicate.

Mixed-effects analysis (historical, unverified). The archived summary fits task random intercepts to 24 task/model combinations from eight tasks and reports $\beta_1 = 1.64$, Wald 95% CI [0.70, 2.58], $p = 0.0006$. With eight clusters, the adequacy of the asymptotic interval/test is uncertain. A previously suggested Kenward-Roger p range was not computed and is withdrawn. The confidence interval and p -value share assumptions; one cannot claim the interval excludes zero under an uncomputed alternative correction. The original fitting code, model diagnostics and task/model data are missing. Random intercepts alone also do not establish independence of tasks sharing the benchmark.

Table 16: Concentration-Drop Relationship by Model Class.

Model	Method	Conc.-Drop ρ	p
LightGBM	Boosting (GBDT)	0.833	0.010
CatBoost	Boosting (ordered)	0.667	0.071
XGBoost	Boosting (histogram)	0.548	0.16
RandomForest	Bagging	0.30	0.47
Mixed-effects, 3 boosting ($n = 24$)		$\beta_1 = 1.64$	<0.001
Mixed-effects, all 4 ($n = 32$)		$\beta_1 = 1.06$	0.010

Interpretation. The historical model/task summaries show heterogeneous outcomes and attribution profiles. They do not establish why RF or MLP correlations differ, nor that differences between correlations are statistically significant. Architectural explanations involving bagging smoothness or boosting reinforcement are hypotheses, requiring controlled comparisons and common attribution definitions.

Feature attribution depends on the fitted predictor, reference distribution and attribution method. These model-specific summaries do not identify causal vulnerability or justify transferring a risk cutoff between model families.

Neural network extension (historical, unverified). The reported two-layer MLP uses permutation importance rather than TreeSHAP, with 10 seeds for sales tasks and 3 for item tasks. Its concentration ratio is therefore not the same attribution estimand as the tree-model SHAP ratio. The reported $\rho = 0.43$, $p = 0.29$ cannot isolate architecture from attribution-method differences. High-drop low-concentration examples motivate a possible multi-feature mechanism but do not prove it; original code/results are required to assess the claim.

K.1 MLP VS LIGHTGBM TASK-LEVEL COMPARISON

Table 17 preserves historical model-family summaries. MLP concentration uses permutation importance, while LightGBM uses SHAP; they are not a common attribution estimand.

Table 17: Task-Level Comparison: MLP vs LightGBM. MLP_C = historical MLP permutation-importance concentration; MLP_Drop = MLP coverage drop (pp); LGB_C = LightGBM SHAP concentration; LGB_Drop = LightGBM coverage drop (pp). Original numerical runs are unverified.

Task	MLP_C	MLP_Drop	LGB_C	LGB_Drop
s-group	27.5%	78.4pp	47.3%	71.2pp
s-payterms	28.0%	60.6pp	54.2%	77.1pp
i-shippoint	77.2%	82.3pp	48.8%	18.5pp
s-incoterms	30.3%	42.3pp	23.7%	8.5pp
s-shipcond	53.1%	67.1pp	50.7%	71.6pp
i-incoterms	31.0%	24.1pp	28.9%	11.3pp
i-plant	28.2%	26.7pp	23.9%	10.6pp
s-office	0.0%	0.1pp	42.6%	0.1pp

These summaries do not distinguish a concentrated causal failure mechanism from a distributed one. In particular, low attribution concentration is not evidence of robustness, and the recovered full SALT panel does not reproduce the historical positive association.

K.2 FAILURE-MECHANISM HYPOTHESES

The original appendix defined concentration-based failure categories and recommended standard deployment below a 40% threshold for boosted trees. We withdraw those definitions as operational criteria and remove the decision table. Neither a thresholded attribution statistic nor a Gini/top-three summary establishes a causal failure mechanism or a coverage guarantee. Concentrated versus distributed sensitivity remain hypotheses requiring controlled feature interventions on a fixed, correctly joined input cohort. The counterexamples and non-replicating follow-up panel rule out interpreting these numerical cutoffs as validated diagnostic boundaries.

K.3 CANDIDATE SENSITIVITY ANALYSES

The following analyses may investigate predictor sensitivity but have not been validated as failure predictors for these data (Table 18).

Table 18: Alternative Diagnostics for Global-Sensitivity Failures in Neural Networks.

Diagnostic	Description
Gradient sensitivity	$S_j = \mathbb{E}[\ \partial f / \partial x_j\ _2]$: average gradient magnitude per feature
Feature ablation	$\Delta_j = \text{Cov}(D) - \text{Cov}(D x_j = \mu_j)$: coverage change when ablating feature j
Score-margin analysis	A coverage transfer bound requires control of score mass near the threshold and a joint-distribution coupling; predictor Lipschitz continuity alone does not make the discontinuous coverage indicator Lipschitz

These are proposed research measurements. They do not identify the cause of historical failures, establish target coverage, or supply a deployment recommendation.

L HYPERPARAMETER SENSITIVITY ANALYSIS

The original appendix reported four configurations, perfect F1 and a mean correlation 0.821 ± 0.012 . The recovered source `compute_hp_sensitivity_real_salt.py` instead loads baseline aggregate concentrations, multiplies them by configuration-specific factors and adds Gaussian noise; it does not refit LightGBM. A separate `compute_hp_sensitivity_full.py` uses synthetic task generators. Neither establishes the claimed four-configuration retraining of real SALT data. No original fitted-model provenance for those tables was recovered. We therefore withdraw their empirical interpretation and remove the unsupported numerical tables from the active revision. They remain in the immutable submitted source. The separately recovered three-seed panel uses one fixed configuration and does not repair this missing hyperparameter experiment.

M SHIFT CHARACTERIZATION

The historical summaries in Table 19 describe marginal covariate and label differences. They do not decompose distribution shift into covariate versus concept mechanisms: changes in $P(X)$, $P(Y)$ or accuracy do not by themselves identify changes in $P(Y | X)$.

Table 19: Shift Characterization for SALT Tasks. Cov KS = mean Kolmogorov–Smirnov statistic across features (train vs. test); %Sig = fraction of features with KS $p < 0.05$; Label JS = Jensen–Shannon divergence of label distributions; Acc Drop = test accuracy minus validation accuracy.

Task	Cov KS	%Sig	Label JS	Acc Drop
s-shipcond	0.50	50%	0.083	+11.9%
s-group	0.50	50%	0.330	+1.0%
s-payterms	0.50	50%	0.078	+4.0%
i-plant	0.86	86%	0.144	−8.7%
i-shippoint	0.86	86%	0.152	+13.1%
s-incoterms	0.50	50%	0.085	+7.6%
i-incoterms	0.86	86%	0.140	+63.4%
s-office	0.50	50%	0.026	−0.2%

The former mechanism-classification column is withdrawn: these marginal summaries do not identify concept shift. Accuracy can change solely because the covariate distribution changes, with $P(Y | X)$ fixed. Historical values remain unverified without original outputs.

Interpretation. The reported covariate KS and label JS values suggest marginal changes, subject to missing provenance and unresolved preprocessing. They do not establish which tasks have concept shift or why coverage changes. In particular, i-incoterms has a reported mean drop of 11.3 pp and should not be classified as catastrophic. Associations between accuracy, concentration and coverage require a common validated pipeline; no conditional-distribution or causal-mechanism conclusion follows from this table alone.

M.1 PER-FEATURE COVARIATE SHIFT DETAILS

Tables 20 and 21 show per-feature KS statistics for sales-level and item-level tasks respectively.

Table 20: Per-Feature Covariate Shift for Sales-Level Tasks (6 features). KS = Kolmogorov–Smirnov statistic (train vs. test).

Feature	Shifted	KS
SALESDOCUMENTTYPE	No	0.0
SALESORGANIZATION	Yes	1.0
DISTRIBUTIONCHANNEL	Yes	1.0
ORGANIZATIONDIVISION	Yes	1.0
BILLINGCOMPANYCODE	No	0.0
TRANSACTIONCURRENCY	No	0.0

For the encoded empirical samples, KS = 1 denotes complete separation. The historical summary reports 3/6 features (50%) as shifted. Original feature-level outputs and a correctly joined input cohort are needed to verify these values and interpret their source.

N HISTORICAL WILDS BENCHMARK SUMMARIES

The original manuscript reported four additional evaluations using WILDS [Koh et al., 2021] and standard datasets. These historical summaries report low concentration and small coverage drops. Candidate generating code has since been recovered, but original outputs and execution provenance remain unavailable, and the recovered protocols conflict with several manuscript descriptions. Covertypes is reused under another split, and CivilComments/Adult are binary tasks outside the selected primary subset. These reports do not establish independent specificity.

Dataset details.

- **CivilComments** (WILDS): Toxicity classification with demographic identity-based shift. TF-IDF features (1,000 dimensions). The historical summary reports concentration 26.2% and maintained coverage; recovered source uses 200 features,

Table 21: Per-Feature Covariate Shift for Item-Level Tasks (7 features). KS = Kolmogorov–Smirnov statistic (train vs. test).

Feature	Shifted	KS
SALESDOCUMENTITEM	Yes	1.0
SALESDOCUMENTITEMCATEGORY	No	0.0
PRODUCT	Yes	1.0
SOLDTOPARTY	Yes	1.0
SHIPTOPARTY	Yes	1.0
BILLTOPARTY	Yes	1.0
PAYERPARTY	Yes	1.0

The historical summary reports 6/7 features (86%) as shifted, consistent with Table 19. Original feature-level outputs are unavailable. The recovered default-join defect prevents attributing these reported differences to pandemic-related changes in products, parties or identifiers.

Table 22: Historical WILDS and Standard Benchmark Summaries (Original Runs Unverified). Two evaluations use WILDS; two use standard datasets. Covertypes also appears in the primary analysis under a different split. All datasets exhibit low SHAP concentration (**all** <40%; mean = 17.0%) and robust coverage, reporting four true negatives among this small selected sample; independent specificity is not established. Shift types: Demographic = identity-based subgroup shift; User/Time = user and temporal domain shift; Geographic = regional covariates; Age = demographic age-based shift. These historical classifications do not establish specificity; for the catastrophic Covertypes case (high- $C \rightarrow$ failure), see Table 11.

Dataset	Source	Shift Type	Conc. (%)	Drop (pp)	Predicted	Actual
CivilComments	WILDS	Demographic	26.2	−0.7	Robust	Robust
Amazon	WILDS	User/Time	1.9	+0.4	Robust	Robust
Covertypes [†]	sklearn	Geographic	12.5	+8.0	Robust	Robust
Adult	OpenML	Age-based	27.5	+1.4	Robust	Robust

[†]The original text attributed differing Covertypes values (49.8% concentration/81.8 pp drop versus 12.5%/8.0 pp) to different wilderness splits. Recovered source for the latter evaluation instead uses a positional split. Without original prediction and split ledgers, neither the protocol identity nor the cause of the difference is established; these two reported points do not validate a concentration-based deployment rule.

not the stated 1,000.

- **Amazon Reviews** (WILDS): Product rating prediction with user/time domain shift. TF-IDF features (1,000 dimensions). The historical summary reports concentration 1.9% and a small drop; original predictions are unavailable.
- **Covertypes** (sklearn): Forest cover classification with geographic shift (wilderness area 1 vs. others). 54 cartographic features. The historical summary reports concentration 12.5% and a small drop; recovered source uses a positional split, conflicting with the stated wilderness split.
- **Adult Income** (OpenML): Income prediction with age-based demographic shift (train: age<50, test: age≥50). 14 demographic/employment features. The historical summary reports concentration 27.5% and maintained coverage; recovered source uses age 40, conflicting with the stated age 50.

Historical classification arithmetic. The four reported benchmark evaluations have low SHAP concentration—all below the 40% threshold (individual values: 26.2%, 1.9%, 12.5%, 27.5%; **mean concentration = 17.0%**). Thresholding the reported values gives four historical true-negative classifications. Selection uncertainty, inconsistent execution protocols and missing raw predictions prevent interpreting this count as validated specificity. The recovered complete SALT panel fails to reproduce the original positive association, so these summaries do not establish a deployable sensitivity/specificity tradeoff.

Combined historical classification arithmetic (SALT + additional benchmarks). Combining the 8 SALT tasks with these 4 additional benchmark summaries yields $n = 12$ total tasks (distinct from the $n = 16$ primary endpoint which uses 8 different external UCI datasets). The threshold accuracy is **91.7% (11/12)**:

- **Sensitivity** (SALT): 4/4 at-risk SALT tasks flagged under the >15 pp rule (historical descriptive counts)
- **Specificity** (additional benchmarks + robust SALT): 7/8 robust tasks correctly unflagged (87.5%)
- **Single misclassification**: sales-office (concentration 42.6%, protective secondary feature)

These small, selected historical counts do not validate an operational boundary. Missing runs, binary tasks outside the primary scope and reuse of Covertypes require explicit treatment in a prospective evaluation.